COMS 4705 – Natural Language Processing – Spring 2015

Assignment 3
Word Sense Disambiguation

(March 13, 2015)

Due Friday, April 3, 11:59 PM

Introduction

What is Word Sense Disambiguation (WSD)?

As we learned in class, word sense disambiguation (WSD) is the task of determining which sense of an ambiguous word is being used in a particular context. The solution to this problem impacts other NLP-related problems such as machine translation and document retrieval.

What is lexical sample task?

The standard WSD (SENSEVAL) task has two variants: "lexical sample" and "all words". The former comprises disambiguating the occurrences of a small sample of target words which were previously selected, while in the latter all the words in a piece of running text need to be disambiguated. In this assignment we will be working on the lexical sample task.

Algorithms for WSD:

There are several types of approaches to WSD, including dictionary-based methods, semi-supervised methods, supervised methods and unsupervised methods. Supervised methods based on sense-labeled corpora are the best-performing methods for sense disambiguation. But such labeled data is expensive and limited. In contrast, dictionary-based methods and unsupervised methods require no labeled texts and are more efficient, but have lower accuracy. In the following part, you will see how two algorithms, a supervised algorithm based on the vector space model and the Lesk algorithm, work on WSD.

What is vector space model?

The Vector space model is a widely-used model in NLP. The basic idea is to represent each instance of an ambiguous word as a vector whose features (attributes) are extracted according to some rules. Usually classification or clustering methods are then applied to solve the problem.

The supervised approach to sense disambiguation can be based on the vector space model.

Here is a skeleton supervised algorithm for the lexical sample task.

---

1. For each instance w_i of word w in a corpus, compute a context vector $\vec{c}$. Label the class of each vector by the sense of word w in the context.
2. Train a **classifier** with these labeled context vectors $\vec{c}$.
3. Disambiguate a particular token t of w by **predicting** the class of token t with the trained classifier.

Later in the assignment we will explain how to extract context vectors and what classification method you can use in order to implement this algorithm.

What is the Lesk algorithm?

The Lesk algorithm is a well-studied dictionary-based algorithm for word sense disambiguation. It is based on the hypothesis that words used together in text are related to each other and that the relation can be observed in the definitions of the words and their senses. Two (or more) words are disambiguated by finding the pair of dictionary senses with the greatest word overlap in their dictionary definitions.

The pseudo code of Lesk is shown below.

```
function SIMPLIFIED LESK(word, sentence) returns best sense of word
    best-sense ← most frequent sense for word
    max-overlap ← 0
    context ← set of words in sentence
    for each sense in senses of word do
        signature ← set of words in the gloss and examples of sense
        overlap ← COMPUTEOVERLAP(signature, context)
        if overlap > max-overlap then
            max-overlap ← overlap
            best-sense ← sense
        end
        return(best-sense)
```

You do not need to implement the Lesk algorithm in this assignment. But as an example you should get a sense of what an unsupervised algorithm is like.

WordNet

WordNet is a lexical database (originally developed for English) that contains information about words, their senses, and the relationships between them. A word with multiple senses is called polysemous. A sense with multiple words is called a "synset" (synonym set). Synsets may be related to each other in dfferent ways, e.g., "cat" is a hypernym (more general concept) of "lion".

Getting Started

How to use WordNet in NLTK for WSD tasks?

- Import wordnet
```
>>> from nltk.corpus import wordnet as wn
```

- Get a word's synset
A synset is a set of words with the same sense. It is identified with a 3-part name of the form: word.pos.nn. We could see the function and examples below:

```
def synsets(self, lemma, pos=None, lang='en'):
```
- Load all synsets with a given lemma (word) and part of speech tag.
- If no pos is specified, all synsets for all parts of speech will be loaded.
- If lang is specified, all the synsets in that language associated with the lemma will be returned. The default value of lang is 'en'.

```
>>> wn.synsets('dog')
[Synset('dog.n.01'), Synset('frump.n.01'), Synset('dog.n.03'), Synset('cad.n.01'), ...]
>>> wn.synsets('dog', pos='v', lang = 'en')
[Synset('chase.v.01')]
```

- Get possible definitions for a word from its synsets
```
>>> for ss in wn.synsets('bank'):
>>> print(ss, ss.definition())
Synset('bank.n.01') sloping land (especially the slope beside a body of water)
Synset('depository_financial_institution.n.01') a financial institution that accepts deposits and channels the money into lending activities
Synset('bank.n.03') a long ridge or pile
Synset('bank.n.04') an arrangement of similar objects in a row or in tiers
...
```

- What languages could we use in NLTK?
The WordNet corpus reader gives access to the Open Multilingual WordNet, using ISO-639 language codes. For example, "cat" stands for Catalan (a language spoken in Spain, France and Andorra, which is related to Spanish, Italian, and the other Romance languages).
```
>>> sorted(wn.langs())
['als', 'arb', 'cat', 'cmn', 'dan', 'eng', 'eus', 'fas','fin', 'fra', 'fre', 'glg', 'heb', 'ind', 'ita', 'jpn', 'nno','nob', 'pol', 'por', 'spa', 'tha', 'zsm']
```

For our lexical sample task, assume we have a Lesk Algorithm implementation in the following format:

lesk(context_sentence, ambiguous_word):
- param str context_sentence: The context sentence where the ambiguous word occurs.
- param str ambiguous_word: The ambiguous word that requires WSD.

- return: "lesk_sense" The Synset() object with the highest signature overlaps.

Usage example:

```
>>> sent = word_tokenize("I went to the bank to deposit money.")
>>> word = "bank"
>>> lesk(sent, word)
```

How to use scikit-learn for machine learning tasks?

- What is scikit-learn?
scikit-learn (http://scikit-learn.org/stable/index.html) is an open source machine learning library for Python. It features various classification, regression and clustering algorithms including support vector machines, k-nearest neighbors, logistic regression, naive Bayes, random forests, gradient boosting, k-means and DBSCAN. For this class, we don't need to understand all the details of these algorithms.

- Loading an example dataset
scikit-learn comes with a few standard datasets. A dataset is a dictionary-like object that holds all the data and some metadata about the data. The data is stored in the .data member. In the case of supervised problems, one or more response variables are stored in the .target member. To load an example dataset,
```
>>> from sklearn import datasets
>>> iris = datasets.load_iris()
>>> digits = datasets.load_digits()
```

- Learning and predicting
Classification is a basic task in machine learning. The task is, given a set of data, usually represented as vectors, and several classes, predict which class each vector belongs to. To solve the problem, we will first train a classifier (estimator) using some data whose classes have already been given. Then we will predict the classes of unlabeled data with the classifier we just trained.

In scikit-learn, an estimator for classification is a Python object that implements the methods **fit(X,y)** and **predict(T)**. An example of an estimator is the class **sklearn.svm.SVC** that implements support vector classification.
```
>>> from sklearn import svm
>>> clf = svm.SVC(gamma=0.001, C=100.)
```

To train the model with our traning dataset
```
>>> clf.fit(digits.data[:-1], digits.target[:-1])
```

Now you can predict new values by
```
>>> clf.predict(digits.data[-1])
```

Here is another example of how to use nearest neighbors algorithm to do the classification.
```
>>> from sklearn import neighbors
>>> clf = neighbors.KNeighborsClassifier(15, weights='uniform')
>>> clf.fit(iris.data, iris.target)
>>> clf.predict(iris.data)
```

Environment variable setup

You can access the provided code by running the following commands:

```
cd ~/hidden/$HIDDENNUMBER/
cp -r ~coms4705/Homework3 .
```

Before running code for this assignment, make sure to run this command:
```
source envsetup.sh
```
to set your PYTHONPATH correctly.

Provided data

We will be using data from the Senseval 3 Lexical Sample Task on multilingual WSD, specifically the English, Catalan, and Spanish corpora.

Datasets:

All the data needed for this assignment are in this directory:
`Homework3/data/`

The data include:

`English-train.xml`	Training data for English
`English-dev.xml`	Development data for English
`English-dev.key`	Answer for the development data for English
`English.sensemap`	Sense map file needed for evaluation
`Catalan-train.xml`	Training data for Catalan
`Catalan-dev.xml`	Development data for Catalan
`Catalan-dev.key`	Answer for the development data for Catalan
`Spanish-train.xml`	Training data for Spanish
`Spanish-dev.xml`	Development data for Spanish
`Spanish-dev.key`	Answer for the development data for Spanish

The data files are in .xml form, where the tags are:

1. **lexelt/lemma**: Each lexelt/lemma represents a word to be disambiguated. The item attribute indicates the word and its POS tag, one lexelt could have several instances.

`<lexelt item="difference.n">`

2. **instance**: Each instance is a case where the target word appears in a context. An instance contains one context.

`<instance id=" difference.n.bnc.00001061">`

3. **context**: The context of the word from each instance. You can choose to use the entire

context or just some words from the context (e.g., common collocations, words within a small window, words with high PMI values, etc).

<context>
In 1991/92 we shall need support even more . Every donation does help . That help makes all the <head>difference</head> to people sick with AIDS who want to stay at home , rather than spend time unnecessarily in hospital . Please help ! SIR JOHN FORD KCMG MG CHAIRMAN OF TRUSTEES
</context>

4. **answer**: The sense of the word in this instance. It only appears in training data.

<answer instance="difference.n.bnc.00001061" senseid="difference%1:24:00::"/>

Note: For Catalan and Spanish data, you will also see <previous>, <target>, <following> tags in a <context> tag. For this assignment we only use the text in the <target> tag.

Note2: If an instance contains multiple <answer> tags, we will only use the first tag for the purpose of training.

You can use *xml* package in Python to parse the .xml file.

WSD systems are evaluated by computing the percentage of ambiguous words for which the WSD system has predicted the correct sense. We have provided an evaluator for you. You will use it to evaluate the performance of the algorithms you implement.

The .key files and .sensemap files are only used for evaluation. You do not need to know the meaning of these files. Later you will see how to use the evaluator with these files.

Assignment

For this assignment, you will be implementing a supervised algorithm based on the vector space model on the Senseval data set. You will need to compare the performance of different features and classifiers, find the best combination, and describe your work in the report.

Provided code

```
baseline.py
scorer2.c
```

Before you start

You will be implementing a supervised algorithm based on the vector space model according to the description in the **Introduction** part. You will need to work on **English, Catalan and Spanish**.

You will disambiguate the target word in each test instance in the development data set. The output should be a single file corresponding to *Language*-dev.xml, with each line for a test instance in the format below:

<div align="center">lexelt_item instance_id sense_id</div>

Note: The lines in the output file have to be sorted in alphabetical ascending order (A-Z). By default, the instances in English-dev.xml **are already in order. Make sure your code keeps this order. Please remove the accent of characters in the output file.**

As a baseline, we explain the sense of a word in any context as the most frequent sense of this word. You can see the output of the baseline by running

<div align="center">python baseline.py *Language*</div>

The output file is *Language*.baseline, with the same format as described above. Your implementation should get a better performance than the baseline. You may be graded based on how good your implementation is.

Implementation

For each language, you should follow the steps below.

1. Compute context vectors for each instance of a word. You will use *Language*-train.xml as training data.

   Suppose in a test instance T_i, the word you need to disambiguate is w. At the beginning you need to tokenize the sentences in the context with nltk.word_tokenize(). S_i is the set of words that are within k distance of word w. Let S be the union of set S_i of size n which may contain duplicate words. Then each test instance will correspond to a vector, where each attribute is the count of a word in set S. For this assignment, we set the window size $k=10$.

2. Train a classifier using the context vectors you obtained. You can use *scikit-learn* library to conduct this step. Here you need to try two different classifiers, k-nearest neighbors (**neighbors.KNeighborsClassifier** class) and SVM (**svm.LinearSVC** class). Compare the performance of the two classifiers in your report.

   You will notice that there are several parameters for the classifiers. For this assignment you can just use the default settings.

3. Use your classifier to perform disambiguation for each test instance in `Language-dev.xml`. The output format should be exactly the same as described above.

4. Try extracting better features than just taking all the words in the window and then redo the classification. You should try the following approaches:

   a) Remove stop words, do stemming, etc.

   b) Add more features by obtaining the synonyms, hyponyms and hypernyms of a word in WordNet.

   c) Use the method described below.

      Suppose the word *w* that needs disambiguation has *k* senses. For a sense *s* of *w* and a word *c* in the window of *w*, we compute the "relevance score" of c with s using

      $$\log\left(\frac{p(s|c)}{p(\bar{s}|c)}\right)$$

      where $p(s|c)$ is the probability that the word *w* has sense *s* when *c* appears, and is computed using

      $$p(s|c) = \frac{N_{s,c}}{N_c}$$

      where $N_{s,c}$ is the number of test instances where the context contains *c* and the sense of word *w* is *s*, and N_c is the number of test instances where the context contains *c*. $p(\bar{s}|c)$ is the probability that the word *w* has sense other than *s* when *c* appears and can be computed similarly.

      For each sense *s*, we compute how relevant the words in the window are, and select the top words as features. The final set of features is the union of all the top features for each sense.

      For example, in `English-train.xml`, the word **activate** has 3 senses, **38201**, **38202** and **38203**. We now compute the relevance score for the word **protein**. There are 10 test instances where the context contains **protein** and the sense is **38201**. There are 8 test instances for **38202** and 1 for **38203**. So the score for (**38201**, **protein**) is $\log\frac{10/19}{9/19}$. (**38202, protein**) gets $\log\frac{8/19}{11/19}$ and (**38203, protein**) gets $\log\frac{1/19}{18/19}$.

   d) Try another good feature extracting method. Hint: *Chi-square or pointwise mutual information* may be useful for extracting features.

   Describe how much each of the approaches above improves the performance.

5. Find the best combination of the feature extracting approaches. And also use the classifier that gives better results. <u>Submit your code with the best performance and the output files named by</u> *Language*<u>.answer</u>.

6. You are welcome to adjust the window size *k* and the classifier parameters to get better results. If you use your own parameters, include in your report how you obtained the settings. You can also try other classifiers if you want. We will compare the performance of all the students' classifiers, and give **bonus points** for the top performers.

Evaluation

You will evaluate the result you obtained from the previous part next. Use scorer2 as shown below to evaluate your output.

scorer2 answer_file key_file [sense_map_file]

where answer_file is the output of your algorithm and key_file is the gold standard for the test data.
Here is an example:

scorer2 English.answer English-dev.key English.sensemap

Note: For Catalan and Spanish data, there is no sensemap file. So you do not need to include it in the command.

Note2: scorer2 is written in C and the binary we give is compiled on the CLIC machines. We have included the source code here in case you want to try it on other platforms.

<u>Include the evaluation results in your report. Include any interesting findings and explain why they occurred</u> (up to 2-3 paragraphs and a table with results).

Submission

As usual, you need to create a directory named Homework3 in your hidden directory and copy all the files into it. The files you should submit are:

main.py
English.answer
Spanish.answer
Catalan.answer
report.txt

Make sure your directory permissions are correct:

chmod 711 ~/hidden
chmod -R 755 ~/hidden/<pin>

Note: You can have multiple source code files, but make sure we can run your algorithm simply by running main.py <input_file> <out_putfile> <language>.